

A1_3

March 27, 2024

0.1 MNIST Data Download

```
[1]: # Load and preprocess MNIST dataset using NumPy  
from keras.datasets import mnist
```

0.2 Part I

```
[5]: import numpy as np  
  
(X_tr, y_tr), (X_t, y_t) = mnist.load_data()  
X_tr = X_tr.reshape(-1, 28*28) / 255.  
X_t = X_t.reshape(-1, 28*28) / 255.  
  
# Define neural network architecture  
class NeuralNetwork:  
    def __init__(self, in_s, hid_s, out_s):  
        self.W1 = np.random.randn(in_s, hid_s)  
        self.b1 = np.zeros(hid_s)  
        self.W2 = np.random.randn(hid_s, out_s)  
        self.b2 = np.zeros(out_s)  
  
    def sigmoid(self, x):  
        return 1 / (1 + np.exp(-x))  
  
    def forward(self, x):  
        self.z1 = np.dot(x, self.W1) + self.b1  
        self.a1 = self.sigmoid(self.z1)  
        self.z2 = np.dot(self.a1, self.W2) + self.b2  
        self.a2 = self.sigmoid(self.z2)  
        return self.a2  
  
# Convert data to NumPy arrays  
in_s = X_tr.shape[1]  
hid_s = 64  
out_s = 10  
X_tr, y_tr = np.array(X_tr), np.array(y_tr)  
X_t, y_t = np.array(X_t), np.array(y_t)
```

```

# Initialize the neural network
model = NeuralNetwork(in_s, hid_s, out_s)

# Define loss function
def mse_loss(predictions, targets):
    return np.mean((predictions - targets)**2)

# Define sigmoid derivative
def sigmoid_derivative(x):
    return x * (1 - x)

# Define learning rate
learning_rate = 0.1

# Train the model
num_epochs = 10
batch_size = 64
num_batches = X_tr.shape[0]

# Train the model with SGD
for epoch in range(num_epochs):
    epoch_l = 0.0
    for i in range(X_tr.shape[0]):
        # Select a random sample
        index = np.random.randint(X_tr.shape[0])
        input_sample = X_tr[index:index+1]
        label_sample = y_tr[index:index+1]

        # Forward pass
        output = model.forward(input_sample)

        # One-hot encode labels
        label_onehot = np.eye(out_s)[label_sample.astype(int)]

        # Compute loss
        loss = mse_loss(output, label_onehot)
        epoch_l += loss

        # Backpropagation
        d_loss = 2 * (output - label_onehot)
        d_z2 = d_loss * sigmoid_derivative(model.a2)
        d_a1 = np.dot(d_z2, model.W2.T)
        d_z1 = d_a1 * sigmoid_derivative(model.a1)

        # Update parameters
        model.W2 -= learning_rate * np.dot(model.a1.T, d_z2)
        model.b2 -= learning_rate * np.sum(d_z2, axis=0)

```

```

        model.W1 -= learning_rate * np.dot(input_sample.T, d_z1)
        model.b1 -= learning_rate * np.sum(d_z1, axis=0)

    epoch_1 /= X_tr.shape[0]
    print(f"Ep [{epoch+1}/{num_epochs}], L: {epoch_1:.4f}")

# Evaluate the model
correct = 0
total = 0

for i in range(0, len(X_t), batch_size):
    inputs = X_t[i:i+batch_size]
    labels = y_t[i:i+batch_size]
    outputs = model.forward(inputs)
    predicted = np.argmax(outputs, axis=1)
    total += batch_size
    correct += np.sum(predicted == labels)

accuracy = correct / total * 100
print(f"Accuracy on the test set: {accuracy:.2f}%")

```

```

Ep [1/10], L: 0.0468
Ep [2/10], L: 0.0197
Ep [3/10], L: 0.0111
Ep [4/10], L: 0.0091
Ep [5/10], L: 0.0082
Ep [6/10], L: 0.0074
Ep [7/10], L: 0.0072
Ep [8/10], L: 0.0066
Ep [9/10], L: 0.0059
Ep [10/10], L: 0.0058
Accuracy on the test set: 94.69%

```

Using the above trained model to predict this image:

```

[7]: from PIL import Image

img = Image.open('img_1.jpg')

# Assuming X_custom contains your own data
X_custom = np.array(img)
X_custom = X_custom.reshape(-1, 28*28)

# Normalize your custom data (if necessary)

```

```

X_custom_normalized = X_custom / 255.0

# Create an empty array to store predictions
predictions = []

# Loop over each sample in the custom data
for sample in X_custom_normalized:
    # Reshape the sample to match the input size of the model
    sample = sample.reshape(1, -1)

    # Forward pass through the model
    output = model.forward(sample)

    # Get the predicted label (index of the maximum value in the output)
    predicted_label = np.argmax(output)

    # Append the predicted label to the list of predictions
    predictions.append(predicted_label)

# Convert the list of predictions to a NumPy array
predictions = np.array(predictions)

# Print the predictions
print("Predictions:", predictions)

```

Predictions: [2]

0.3 Pytorch Dependencies

```

[21]: # Setting up Pytorch imports and Data

import torch
import torch.nn as nn
import torch.optim as opt
import torchvision.datasets as dataset
import torchvision.transforms as trans
from torch.utils.data import DataLoader

# Set device
dev = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Load MNIST dataset and preprocess
trans = trans.Compose([
    trans.ToTensor(),
    trans.Normalize((0.5,), (0.5,))
])

```

```

tr_dataset = dataset.MNIST(root='./data', train=True, transform=trans,
    ↳download=True)
t_dataset = dataset.MNIST(root='./data', train=False, transform=trans,
    ↳download=True)

tr_loader = DataLoader(tr_dataset, batch_size=batch_size, shuffle=True)
t_loader = DataLoader(t_dataset, batch_size=batch_size, shuffle=False)

```

0.4 Part II

```

[25]: # Define the autoencoder architecture
class Autoencoder(nn.Module):
    def __init__(self, in_s, hid_s):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(in_s, hid_s),
            nn.Tanh()
        )
        self.decoder = nn.Linear(hid_s, in_s)

    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x

# Initialize the autoencoder
in_s = 28 * 28 # MNIST image size
hid_s = 64
autoenc = Autoencoder(in_s, hid_s).to(dev)

# Loss and Optimizer
c = nn.MSELoss()
op = opt.Adam(autoenc.parameters(), lr=0.001)

# Train the autoencoder
num_epochs = 10
for ep in range(num_epochs):
    r_loss = 0.0
    for img, _ in tr_loader:
        img = img.view(img.size(0), -1).to(dev)

        out = autoenc(img)

        loss = c(out, img)

        op.zero_grad()
        loss.backward()

```

```

        op.step()

        r_loss += loss.item()

    ep_loss = r_loss / len(tr_loader)
    print(f"Ep [{ep+1}/{num_epochs}], L: {ep_loss:.4f}")

```

```

Ep [1/10], L: 0.1265
Ep [2/10], L: 0.0585
Ep [3/10], L: 0.0505
Ep [4/10], L: 0.0470
Ep [5/10], L: 0.0455
Ep [6/10], L: 0.0442
Ep [7/10], L: 0.0433
Ep [8/10], L: 0.0428
Ep [9/10], L: 0.0423
Ep [10/10], L: 0.0420

```

0.5 Part III

```

[23]: # Define the neural network architecture
class MLP(nn.Module):
    def __init__(self, in_s, hid_s1, hid_s2, out_s):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(in_s, hid_s1)
        self.fc2 = nn.Linear(hid_s1, hid_s2)
        self.fc3 = nn.Linear(hid_s2, out_s)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.sigmoid(self.fc1(x))
        x = self.sigmoid(self.fc2(x))
        x = self.fc3(x)
        return x

# Hyperparameters
in_s = 784 # 28x28 for MNIST images
hid_s1 = 256
hid_s2 = 128
out_s = 10 # 10 classes for MNIST digits
lr = 0.2
batch_size = 64
num_epochs = 10

# Initialize the model

```

```

model = MLP(in_s, hid_s1, hid_s2, out_s)

# Training loop
tot_stp = len(tr_loader)
for ep in range(num_epochs):
    for i, (img, lbl) in enumerate(tr_loader):
        # Flatten
        img = img.view(-1, 28*28)

        # Forward pass
        out = model(img)

        # Compute MSE loss
        c = nn.MSELoss()
        loss = c(out, torch.nn.functional.one_hot(lbl, num_classes=out_s).
        ↪float())

        # Zero gradients, backward pass, and optimize
        model.zero_grad()
        loss.backward()

        # SGD update
        with torch.no_grad():
            for par in model.parameters():
                par -= lr * par.grad

        if (i+1) % 100 == 0:
            print ('Ep [{}/{}], S [{}/{}], L: {:.4f}'
                    .format(ep+1, num_epochs, i+1, tot_stp, loss.item()))

# Test the model
with torch.no_grad():
    corr = 0
    total = 0
    for img, lbl in t_loader:
        img = img.view(-1, 28*28)
        out = model(img)
        _, predicted = torch.max(out.data, 1)
        total += lbl.size(0)
        corr += (predicted == lbl).sum().item()

    print('Accuracy of the network on the 10000 test images: {:.2f} %'.
    ↪format(100 * corr / total))

```

Ep [1/10], S [100/938], L: 0.0918

Ep [1/10], S [200/938], L: 0.0869

Ep [1/10], S [300/938], L: 0.0881

Ep [1/10], S [400/938], L: 0.0838
 Ep [1/10], S [500/938], L: 0.0789
 Ep [1/10], S [600/938], L: 0.0718
 Ep [1/10], S [700/938], L: 0.0695
 Ep [1/10], S [800/938], L: 0.0708
 Ep [1/10], S [900/938], L: 0.0665
 Ep [2/10], S [100/938], L: 0.0626
 Ep [2/10], S [200/938], L: 0.0567
 Ep [2/10], S [300/938], L: 0.0578
 Ep [2/10], S [400/938], L: 0.0549
 Ep [2/10], S [500/938], L: 0.0550
 Ep [2/10], S [600/938], L: 0.0588
 Ep [2/10], S [700/938], L: 0.0544
 Ep [2/10], S [800/938], L: 0.0523
 Ep [2/10], S [900/938], L: 0.0510
 Ep [3/10], S [100/938], L: 0.0415
 Ep [3/10], S [200/938], L: 0.0541
 Ep [3/10], S [300/938], L: 0.0515
 Ep [3/10], S [400/938], L: 0.0463
 Ep [3/10], S [500/938], L: 0.0426
 Ep [3/10], S [600/938], L: 0.0403
 Ep [3/10], S [700/938], L: 0.0463
 Ep [3/10], S [800/938], L: 0.0475
 Ep [3/10], S [900/938], L: 0.0493
 Ep [4/10], S [100/938], L: 0.0466
 Ep [4/10], S [200/938], L: 0.0467
 Ep [4/10], S [300/938], L: 0.0554
 Ep [4/10], S [400/938], L: 0.0423
 Ep [4/10], S [500/938], L: 0.0446
 Ep [4/10], S [600/938], L: 0.0422
 Ep [4/10], S [700/938], L: 0.0452
 Ep [4/10], S [800/938], L: 0.0361
 Ep [4/10], S [900/938], L: 0.0393
 Ep [5/10], S [100/938], L: 0.0418
 Ep [5/10], S [200/938], L: 0.0411
 Ep [5/10], S [300/938], L: 0.0366
 Ep [5/10], S [400/938], L: 0.0466
 Ep [5/10], S [500/938], L: 0.0390
 Ep [5/10], S [600/938], L: 0.0396
 Ep [5/10], S [700/938], L: 0.0446
 Ep [5/10], S [800/938], L: 0.0406
 Ep [5/10], S [900/938], L: 0.0313
 Ep [6/10], S [100/938], L: 0.0471
 Ep [6/10], S [200/938], L: 0.0364
 Ep [6/10], S [300/938], L: 0.0424
 Ep [6/10], S [400/938], L: 0.0460
 Ep [6/10], S [500/938], L: 0.0351
 Ep [6/10], S [600/938], L: 0.0388


```

Ep [6/10], S [700/938], L: 0.0328
Ep [6/10], S [800/938], L: 0.0389
Ep [7/10], S [100/938], L: 0.0388
Ep [7/10], S [200/938], L: 0.0404
Ep [7/10], S [300/938], L: 0.0331
Ep [7/10], S [400/938], L: 0.0354
Ep [7/10], S [500/938], L: 0.0389
Ep [7/10], S [600/938], L: 0.0317
Ep [7/10], S [700/938], L: 0.0350
Ep [7/10], S [800/938], L: 0.0295
Ep [7/10], S [900/938], L: 0.0309
Ep [8/10], S [100/938], L: 0.0371
Ep [8/10], S [200/938], L: 0.0272
Ep [8/10], S [300/938], L: 0.0298
Ep [8/10], S [400/938], L: 0.0288
Ep [8/10], S [500/938], L: 0.0370
Ep [8/10], S [600/938], L: 0.0299
Ep [8/10], S [700/938], L: 0.0267
Ep [8/10], S [800/938], L: 0.0335
Ep [8/10], S [900/938], L: 0.0319
Ep [9/10], S [100/938], L: 0.0267
Ep [9/10], S [200/938], L: 0.0324
Ep [9/10], S [300/938], L: 0.0333
Ep [9/10], S [400/938], L: 0.0289
Ep [9/10], S [500/938], L: 0.0312
Ep [9/10], S [600/938], L: 0.0288
Ep [9/10], S [700/938], L: 0.0332
Ep [9/10], S [800/938], L: 0.0259
Ep [9/10], S [900/938], L: 0.0273
Ep [10/10], S [100/938], L: 0.0272
Ep [10/10], S [200/938], L: 0.0320
Ep [10/10], S [300/938], L: 0.0292
Ep [10/10], S [400/938], L: 0.0205
Ep [10/10], S [500/938], L: 0.0256
Ep [10/10], S [600/938], L: 0.0247
Ep [10/10], S [700/938], L: 0.0238
Ep [10/10], S [800/938], L: 0.0298
Ep [10/10], S [900/938], L: 0.0392
Accuracy of the network on the 10000 test images: 89.18 %

```

0.6 Part IV

```

[27]: # Define Autoencoder architecture
class Autoencoder(nn.Module):
    def __init__(self, in_s, hid_s):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Linear(in_s, hid_s)

```

```

        self.decoder = nn.Linear(hid_s, in_s)

    def forward(self, x):
        x = torch.sigmoid(self.encoder(x))
        x = torch.sigmoid(self.decoder(x))
        return x

# Hyperparameters
in_s = 784 # 28x28 for MNIST images
hid_s = 256
batch_size = 64
num_epochs = 10

# Initialize the autoencoder
autoenc = Autoencoder(in_s, hid_s)

# Define the loss function and optimizer
c = nn.MSELoss()
op = opt.SGD(autoenc.parameters(), lr=0.1)

# Training loop for the autoencoder
tot_stp = len(tr_loader)
for ep in range(num_epochs):
    for i, (img, _) in enumerate(tr_loader):
        img = img.view(-1, in_s)

        # Forward pass
        out = autoenc(img)
        loss = c(out, img)

        # Backward and optimize
        op.zero_grad()
        loss.backward()
        op.step()

        if (i+1) % 100 == 0:
            print ('Autoencoder: Ep [{}/{}], S [{}/{}], L: {:.4f}'
                  .format(ep+1, num_epochs, i+1, tot_stp, loss.item()))

# Save the pre-trained weights of the first layer
pretrained_weights = autoenc.encoder.weight.detach().clone()

# Define the MLP architecture with pre-trained weights
class MLP(nn.Module):
    def __init__(self, in_s, hid_s1, hid_s2, out_s):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(in_s, hid_s1)

```

```

self.fc2 = nn.Linear(hid_s1, hid_s2)
self.fc3 = nn.Linear(hid_s2, out_s)
self.sigmoid = nn.Sigmoid()

# Initialize first two layers with pre-trained weights
with torch.no_grad():
    self.fc1.weight = nn.Parameter(pretrained_weights)
    self.fc1.requires_grad = False # Freeze the parameters

def forward(self, x):
    x = self.sigmoid(self.fc1(x))
    x = self.sigmoid(self.fc2(x))
    x = self.sigmoid(self.fc3(x))
    return x

# Re-initialize the autoencoder for the fine-tuning phase
autoenc = Autoencoder(in_s, hid_s)

# Load MNIST dataset again for fine-tuning
tr_loader = DataLoader(tr_dataset, batch_size=batch_size, shuffle=True)

# Initialize the MLP model with pre-trained weights
model = MLP(in_s, hid_s1=hid_s, hid_s2=128, out_s=10)

# Define the loss function and optimizer
c = nn.MSELoss()
op = opt.SGD(model.parameters(), lr=0.1)

# Training loop for fine-tuning the entire MLP
for ep in range(num_epochs):
    for i, (img, lbl) in enumerate(tr_loader):
        img = img.view(-1, in_s)

        # Forward pass
        out = model(img)
        lbl = torch.nn.functional.one_hot(lbl, num_classes=10).float()
        loss = c(out, lbl)

        # Backward and optimize
        op.zero_grad()
        loss.backward()
        op.step()

    if (i+1) % 100 == 0:
        print ('Fine-tuning: Ep [{}/{}], S [{}/{}], L: {:.4f}'
              .format(ep+1, num_epochs, i+1, tot_stp, loss.item()))

```

Autoencoder: Ep [1/10], S [100/938], L: 1.4770
 Autoencoder: Ep [1/10], S [200/938], L: 1.1810
 Autoencoder: Ep [1/10], S [300/938], L: 1.0684
 Autoencoder: Ep [1/10], S [400/938], L: 1.0222
 Autoencoder: Ep [1/10], S [500/938], L: 1.0032
 Autoencoder: Ep [1/10], S [600/938], L: 0.9807
 Autoencoder: Ep [1/10], S [700/938], L: 0.9776
 Autoencoder: Ep [1/10], S [800/938], L: 0.9693
 Autoencoder: Ep [1/10], S [900/938], L: 0.9638
 Autoencoder: Ep [2/10], S [100/938], L: 0.9585
 Autoencoder: Ep [2/10], S [200/938], L: 0.9520
 Autoencoder: Ep [2/10], S [300/938], L: 0.9516
 Autoencoder: Ep [2/10], S [400/938], L: 0.9499
 Autoencoder: Ep [2/10], S [500/938], L: 0.9467
 Autoencoder: Ep [2/10], S [600/938], L: 0.9468
 Autoencoder: Ep [2/10], S [700/938], L: 0.9458
 Autoencoder: Ep [2/10], S [800/938], L: 0.9425
 Autoencoder: Ep [2/10], S [900/938], L: 0.9428
 Autoencoder: Ep [3/10], S [100/938], L: 0.9417
 Autoencoder: Ep [3/10], S [200/938], L: 0.9372
 Autoencoder: Ep [3/10], S [300/938], L: 0.9365
 Autoencoder: Ep [3/10], S [400/938], L: 0.9417
 Autoencoder: Ep [3/10], S [500/938], L: 0.9407
 Autoencoder: Ep [3/10], S [600/938], L: 0.9411
 Autoencoder: Ep [3/10], S [700/938], L: 0.9350
 Autoencoder: Ep [3/10], S [800/938], L: 0.9370
 Autoencoder: Ep [3/10], S [900/938], L: 0.9336
 Autoencoder: Ep [4/10], S [100/938], L: 0.9369
 Autoencoder: Ep [4/10], S [200/938], L: 0.9366
 Autoencoder: Ep [4/10], S [300/938], L: 0.9398
 Autoencoder: Ep [4/10], S [400/938], L: 0.9308
 Autoencoder: Ep [4/10], S [500/938], L: 0.9385
 Autoencoder: Ep [4/10], S [600/938], L: 0.9308
 Autoencoder: Ep [4/10], S [700/938], L: 0.9340
 Autoencoder: Ep [4/10], S [800/938], L: 0.9346
 Autoencoder: Ep [4/10], S [900/938], L: 0.9329
 Autoencoder: Ep [5/10], S [100/938], L: 0.9324
 Autoencoder: Ep [5/10], S [200/938], L: 0.9353
 Autoencoder: Ep [5/10], S [300/938], L: 0.9342
 Autoencoder: Ep [5/10], S [400/938], L: 0.9311
 Autoencoder: Ep [5/10], S [500/938], L: 0.9294
 Autoencoder: Ep [5/10], S [600/938], L: 0.9298
 Autoencoder: Ep [5/10], S [700/938], L: 0.9317
 Autoencoder: Ep [5/10], S [800/938], L: 0.9290
 Autoencoder: Ep [5/10], S [900/938], L: 0.9290
 Autoencoder: Ep [6/10], S [100/938], L: 0.9319
 Autoencoder: Ep [6/10], S [200/938], L: 0.9315
 Autoencoder: Ep [6/10], S [300/938], L: 0.9310

Autoencoder: Ep [6/10], S [400/938], L: 0.9297
 Autoencoder: Ep [6/10], S [500/938], L: 0.9310
 Autoencoder: Ep [6/10], S [600/938], L: 0.9328
 Autoencoder: Ep [6/10], S [700/938], L: 0.9308
 Autoencoder: Ep [6/10], S [800/938], L: 0.9326
 Autoencoder: Ep [6/10], S [900/938], L: 0.9309
 Autoencoder: Ep [7/10], S [100/938], L: 0.9334
 Autoencoder: Ep [7/10], S [200/938], L: 0.9297
 Autoencoder: Ep [7/10], S [300/938], L: 0.9309
 Autoencoder: Ep [7/10], S [400/938], L: 0.9309
 Autoencoder: Ep [7/10], S [500/938], L: 0.9315
 Autoencoder: Ep [7/10], S [600/938], L: 0.9298
 Autoencoder: Ep [7/10], S [700/938], L: 0.9312
 Autoencoder: Ep [7/10], S [800/938], L: 0.9284
 Autoencoder: Ep [7/10], S [900/938], L: 0.9305
 Autoencoder: Ep [8/10], S [100/938], L: 0.9277
 Autoencoder: Ep [8/10], S [200/938], L: 0.9302
 Autoencoder: Ep [8/10], S [300/938], L: 0.9297
 Autoencoder: Ep [8/10], S [400/938], L: 0.9286
 Autoencoder: Ep [8/10], S [500/938], L: 0.9317
 Autoencoder: Ep [8/10], S [600/938], L: 0.9272
 Autoencoder: Ep [8/10], S [700/938], L: 0.9315
 Autoencoder: Ep [8/10], S [800/938], L: 0.9279
 Autoencoder: Ep [8/10], S [900/938], L: 0.9291
 Autoencoder: Ep [9/10], S [100/938], L: 0.9291
 Autoencoder: Ep [9/10], S [200/938], L: 0.9296
 Autoencoder: Ep [9/10], S [300/938], L: 0.9294
 Autoencoder: Ep [9/10], S [400/938], L: 0.9290
 Autoencoder: Ep [9/10], S [500/938], L: 0.9243
 Autoencoder: Ep [9/10], S [600/938], L: 0.9298
 Autoencoder: Ep [9/10], S [700/938], L: 0.9292
 Autoencoder: Ep [9/10], S [800/938], L: 0.9306
 Autoencoder: Ep [9/10], S [900/938], L: 0.9288
 Autoencoder: Ep [10/10], S [100/938], L: 0.9299
 Autoencoder: Ep [10/10], S [200/938], L: 0.9291
 Autoencoder: Ep [10/10], S [300/938], L: 0.9302
 Autoencoder: Ep [10/10], S [400/938], L: 0.9309
 Autoencoder: Ep [10/10], S [500/938], L: 0.9292
 Autoencoder: Ep [10/10], S [600/938], L: 0.9291
 Autoencoder: Ep [10/10], S [700/938], L: 0.9292
 Autoencoder: Ep [10/10], S [800/938], L: 0.9304
 Autoencoder: Ep [10/10], S [900/938], L: 0.9285
 Fine-tuning: Ep [1/10], S [100/938], L: 0.0910
 Fine-tuning: Ep [1/10], S [200/938], L: 0.0900
 Fine-tuning: Ep [1/10], S [300/938], L: 0.0898
 Fine-tuning: Ep [1/10], S [400/938], L: 0.0899
 Fine-tuning: Ep [1/10], S [500/938], L: 0.0900
 Fine-tuning: Ep [1/10], S [600/938], L: 0.0901

Fine-tuning: Ep [1/10], S [700/938], L: 0.0897
 Fine-tuning: Ep [1/10], S [800/938], L: 0.0901
 Fine-tuning: Ep [1/10], S [900/938], L: 0.0900
 Fine-tuning: Ep [2/10], S [100/938], L: 0.0902
 Fine-tuning: Ep [2/10], S [200/938], L: 0.0900
 Fine-tuning: Ep [2/10], S [300/938], L: 0.0896
 Fine-tuning: Ep [2/10], S [400/938], L: 0.0899
 Fine-tuning: Ep [2/10], S [500/938], L: 0.0900
 Fine-tuning: Ep [2/10], S [600/938], L: 0.0898
 Fine-tuning: Ep [2/10], S [700/938], L: 0.0900
 Fine-tuning: Ep [2/10], S [800/938], L: 0.0895
 Fine-tuning: Ep [2/10], S [900/938], L: 0.0902
 Fine-tuning: Ep [3/10], S [100/938], L: 0.0900
 Fine-tuning: Ep [3/10], S [200/938], L: 0.0899
 Fine-tuning: Ep [3/10], S [300/938], L: 0.0899
 Fine-tuning: Ep [3/10], S [400/938], L: 0.0898
 Fine-tuning: Ep [3/10], S [500/938], L: 0.0902
 Fine-tuning: Ep [3/10], S [600/938], L: 0.0900
 Fine-tuning: Ep [3/10], S [700/938], L: 0.0899
 Fine-tuning: Ep [3/10], S [800/938], L: 0.0902
 Fine-tuning: Ep [3/10], S [900/938], L: 0.0902
 Fine-tuning: Ep [4/10], S [100/938], L: 0.0899
 Fine-tuning: Ep [4/10], S [200/938], L: 0.0898
 Fine-tuning: Ep [4/10], S [300/938], L: 0.0900
 Fine-tuning: Ep [4/10], S [400/938], L: 0.0899
 Fine-tuning: Ep [4/10], S [500/938], L: 0.0900
 Fine-tuning: Ep [4/10], S [600/938], L: 0.0898
 Fine-tuning: Ep [4/10], S [700/938], L: 0.0899
 Fine-tuning: Ep [4/10], S [800/938], L: 0.0899
 Fine-tuning: Ep [4/10], S [900/938], L: 0.0899
 Fine-tuning: Ep [5/10], S [100/938], L: 0.0901
 Fine-tuning: Ep [5/10], S [200/938], L: 0.0901
 Fine-tuning: Ep [5/10], S [300/938], L: 0.0899
 Fine-tuning: Ep [5/10], S [400/938], L: 0.0902
 Fine-tuning: Ep [5/10], S [500/938], L: 0.0899
 Fine-tuning: Ep [5/10], S [600/938], L: 0.0898
 Fine-tuning: Ep [5/10], S [700/938], L: 0.0899
 Fine-tuning: Ep [5/10], S [800/938], L: 0.0900
 Fine-tuning: Ep [5/10], S [900/938], L: 0.0902
 Fine-tuning: Ep [6/10], S [100/938], L: 0.0901
 Fine-tuning: Ep [6/10], S [200/938], L: 0.0900
 Fine-tuning: Ep [6/10], S [300/938], L: 0.0904
 Fine-tuning: Ep [6/10], S [400/938], L: 0.0901
 Fine-tuning: Ep [6/10], S [500/938], L: 0.0898
 Fine-tuning: Ep [6/10], S [600/938], L: 0.0900
 Fine-tuning: Ep [6/10], S [700/938], L: 0.0901
 Fine-tuning: Ep [6/10], S [800/938], L: 0.0898
 Fine-tuning: Ep [6/10], S [900/938], L: 0.0899

```

Fine-tuning: Ep [7/10], S [100/938], L: 0.0901
Fine-tuning: Ep [7/10], S [200/938], L: 0.0901
Fine-tuning: Ep [7/10], S [300/938], L: 0.0900
Fine-tuning: Ep [7/10], S [400/938], L: 0.0901
Fine-tuning: Ep [7/10], S [500/938], L: 0.0897
Fine-tuning: Ep [7/10], S [600/938], L: 0.0900
Fine-tuning: Ep [7/10], S [700/938], L: 0.0903
Fine-tuning: Ep [7/10], S [800/938], L: 0.0900
Fine-tuning: Ep [7/10], S [900/938], L: 0.0901
Fine-tuning: Ep [8/10], S [100/938], L: 0.0900
Fine-tuning: Ep [8/10], S [200/938], L: 0.0900
Fine-tuning: Ep [8/10], S [300/938], L: 0.0902
Fine-tuning: Ep [8/10], S [400/938], L: 0.0902
Fine-tuning: Ep [8/10], S [500/938], L: 0.0898
Fine-tuning: Ep [8/10], S [600/938], L: 0.0898
Fine-tuning: Ep [8/10], S [700/938], L: 0.0901
Fine-tuning: Ep [8/10], S [800/938], L: 0.0901
Fine-tuning: Ep [8/10], S [900/938], L: 0.0902
Fine-tuning: Ep [9/10], S [100/938], L: 0.0902
Fine-tuning: Ep [9/10], S [200/938], L: 0.0900
Fine-tuning: Ep [9/10], S [300/938], L: 0.0899
Fine-tuning: Ep [9/10], S [400/938], L: 0.0899
Fine-tuning: Ep [9/10], S [500/938], L: 0.0901
Fine-tuning: Ep [9/10], S [600/938], L: 0.0901
Fine-tuning: Ep [9/10], S [700/938], L: 0.0899
Fine-tuning: Ep [9/10], S [800/938], L: 0.0898
Fine-tuning: Ep [9/10], S [900/938], L: 0.0899
Fine-tuning: Ep [10/10], S [100/938], L: 0.0899
Fine-tuning: Ep [10/10], S [200/938], L: 0.0901
Fine-tuning: Ep [10/10], S [300/938], L: 0.0900
Fine-tuning: Ep [10/10], S [400/938], L: 0.0901
Fine-tuning: Ep [10/10], S [500/938], L: 0.0902
Fine-tuning: Ep [10/10], S [600/938], L: 0.0899
Fine-tuning: Ep [10/10], S [700/938], L: 0.0901
Fine-tuning: Ep [10/10], S [800/938], L: 0.0899
Fine-tuning: Ep [10/10], S [900/938], L: 0.0897

```

0.7 Part V

```

[ ]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms

# Define LeNet model
class LeNet(nn.Module):

```

```

def __init__(self):
    super(LeNet, self).__init__()
    self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
    self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
    self.fc1 = nn.Linear(16 * 4 * 4, 120)
    self.fc2 = nn.Linear(120, 84)
    self.fc3 = nn.Linear(84, 10)

def forward(self, x):
    x = torch.nn.functional.relu(self.conv1(x))
    x = torch.nn.functional.max_pool2d(x, kernel_size=2, stride=2)
    x = torch.nn.functional.relu(self.conv2(x))
    x = torch.nn.functional.max_pool2d(x, kernel_size=2, stride=2)
    x = torch.flatten(x, 1)
    x = torch.nn.functional.relu(self.fc1(x))
    x = torch.nn.functional.relu(self.fc2(x))
    x = self.fc3(x)
    return x

# Load MNIST dataset
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.
    ↪5,), (0.5,))])
trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True,
    ↪transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True)

# Initialize LeNet model
net = LeNet()

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

# Training loop
for epoch in range(5): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        inputs, labels = data
        optimizer.zero_grad()
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

    running_loss += loss.item()
    if i % 1000 == 999: # print every 1000 mini-batches

```



```

        print('%d, %5d] loss: %.3f' % (epoch + 1, i + 1, running_loss /
↪1000))
        running_loss = 0.0

print('Finished Training')

```

```

[1, 1000] loss: 1.675
[2, 1000] loss: 0.157
[3, 1000] loss: 0.100
[4, 1000] loss: 0.079
[5, 1000] loss: 0.064
Finished Training

```

```

[ ]: import torch
import torchvision
import torchvision.transforms as transforms

# Load the test dataset
testset = torchvision.datasets.MNIST(root='./data', train=False, download=True,
↪transform=transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
]))

testloader = torch.utils.data.DataLoader(testset, batch_size=32, shuffle=False)

# Function to calculate accuracy
def calculate_accuracy(net, dataloader):
    correct = 0
    total = 0
    with torch.no_grad():
        for data in dataloader:
            images, labels = data
            outputs = net(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    return correct / total

# Calculate accuracy on test dataset
accuracy = calculate_accuracy(net, testloader)
print('Accuracy of the network on the test images: {:.2f}%'.format(accuracy *
↪100))

```

Accuracy of the network on the test images: 98.13%